

Assignment	Write a program that downloads a URL using multiple threads to increase the download speed. Each thread downloads a part of the URL using the http protocol. The application should have a graphical interface where the user can a) input the URL to download, b) view download progress and speed, and c) cancel a pending download.
What to Submit	1. Create a project named flashget and commit it to Github Classroom using the repository that will be created for you. The Github classroom assignment URL is being sent by email. Go to the URL and "accept" the assignment: 2. Create a <i>runnable JAR file</i> of your application with the name flashget.jar in the top level folder of your project.

Requirements

Write a graphical application that downloads a file using threads, where each thread downloads one part of the file and all threads run simultaneously. The default should be 5 threads.

For example, to download a 100KB file using 5 threads, thread 1 would download bytes 0-20479, thread 2 would download bytes 20480-40959, thread 3 would download bytes 40960-61439, etc. Java provides random access I/O, so (for example) thread 2 can *seek* to byte position 1.2M in the output file and start writing there. All 5 threads can write to the same output file, provided they are careful not to write to same location. (See below for details.)

The user interface should display the progress of each thread and the overall download progress. The UI should remain responsive during download, and permit the user to cancel a pending download.

The program should never crash, throw exceptions, or print on the console. When an exception occurs, catch it and display a message in the graphical user interface.

What Do You Need To Know?

To implement a solution to this, you need to know the following

1. How to connect to a URL in Java
2. How to get the *size* of the file at a URL
3. How to download a part of a URL (not the whole file), by specifying a range of bytes
4. How to *write* to an arbitrary location in a local file
5. How to manage several threads as a group

Each of these is described below.

1. How to connect to a URL in Java?

Java has `java.net.URL` and `java.net.URI` classes with useful methods for using URLs, getting information about a URL, and reading a URL. To connect to a URL for reading, use:

```
String name = "http://www.ku.ac.th/index.html";
URL url = null;
try {
    URL url = new URL( name );
} catch ( MalformedURLException e ) {
    System.err.println( e.getMessage() );
    return null;
}
// get in input stream (may throw IOException)
InputStream instream = url.openStream( );
```

2. How to you get the size of the file at a URL?

The `URL` class has a `getConnection()` method that lets you query and manipulate an HTTP connection. You must call this method *before* calling `openStream` if you want to set any parameters on the connection.

The `URLConnection` class has several query methods, including `getContentLengthLong()`. As you can guess, this method returns the length of the file or resource (if known). If you are certain that the file is smaller than 2GB you can also use `getContentLength()` which returns an `int`. If the size of the file or resource is not known, `getContentLength()` and `getContentLengthLong()` will return -1.

Open a URL and query the content length.

```
// use the code above to create the URL object
// but don't call openStream.
URL url = ... // from previous code
try {
    URLConnection connection = url.openConnection( );
    long length = connection.getContentLengthLong( );
} catch (IOException ioe) {
    // handle it
}
```

3. How do you read part of a URL, from a given starting position (a byte index)?

The HTTP 1.1 (and newer) protocol allows the client to specify many *request headers* that affect processing of the HTTP request. One of the headers is named "Range" which requests the range of bytes to get. Here is an example HTTP "GET" request with a **Range** header set to get bytes 1,000 - 1,999. It also has two other headers.

```
GET /path/index.html HTTP/1.1
UserAgent: Mozilla
Host: www.example.com
Range: bytes=1000-1999
```

If you want to read from a given byte to the end of the file or resource, use the header "Range: bytes=1000-" (no ending byte number).

In Java, you set HTTP request headers using `setRequestProperty()` of the `URLConnection` object. Here is an example of setting the byte range 4096 - 8191 (bytes 4K - 8K).

```
int start = 4096;
int size = 4096;
URLConnection connection = url.openConnection();
String range = null;
if ( size > 0 ) {
    range = String.format("bytes=%d-%d", start, start+size-1);
}
else {
    // size not given, so read from start byte to end of file
    range = String.format("bytes=%d-", start);
}
connection.setRequestProperty( "Range", range );

// now get the input stream for reading the part of the URL
InputStream instream = connection.getInputStream();
```

4. How do you write to an arbitrary position in a file?

Suppose you want to start writing to a file at a given location (start) rather than writing from the *beginning* of the file? Java has a Random Access File type that lets you *seek* to any location in a file and then write (or read) at that location. For example:

```
File file = new File( "filename.txt" );
int start = 1000;
int size = 1000;
private RandomAccessFile writer; // for output
writer = new RandomAccessFile( file, "rwd" );
        // see Javadoc for meaning of "rwd" flags
writer.seek( start );
```

You can seek() to a location beyond the end of the file, too. Java will extend the file's size.

After executing the above statements to open a RandomAccessFile and seek to the place where we want to start writing, we can read data from the URLConnection (instream) and write to the output file.

```
// BUFFER_SIZE is a constant for the size of each read.
// Try 4096 and increase it if it helps performance.
int buffsize = Math.min( size, BUFFER_SIZE );
byte [] buffer = new byte[buffsize];
int bytesRead = 0;
try {
    do {
        int n = instream.read( buffer );
        if ( n < 0 ) break;
        writer.write(buffer, 0, n);
        bytesRead += n;
    } while ( bytesRead < size );
} catch (IOException e) { ... }
finally {
    try {
        if ( instream != null ) instream.close();
    } catch (IOException e) { /* ignore it */ }
}
```

Hard disk and SSDs are written to in sectors or blocks at a time (not bytes). For a hard disk the sector size is typically 4,096 bytes. So for efficiency, try to perform reads and writes as multiples of 4096 (instead of, say, 4,000).

In the code above, `instream.read()` returns an int which is actual number of bytes read. This may be smaller than the buffer size, so don't assume that read() fills the entire buffer array.

The `finally` block is to ensure the connection to the URL is closed when it is no longer needed.

For a multi-threaded downloader, each thread will open its own InputStream (using URLConnection) and its own output stream (using RandomAccessFile).

5. Managing Threads as a Group

In your application you may want to know when all the downloader threads have finished, or to cancel them all. There are a few ways to manage a collection of threads. The lowest level approach is to create a ThreadGroup. This is in the case you are creating and running threads yourself. ThreadGroup has an activeCount method that returns the number of active (live) threads in the group, and a method to interrupt all threads. (It is up to the threads to check for interrupts.)

A higher level way to use an *ExecutorService* to create and manage threads.